

Homework Set 2, CPSC 8420, Fall 2023

Parampreet Singh

Due 10/26/2023, Thursday, 11:59PM EST

Problem 1

For PCA, from the perspective of maximizing variance, please show that the solution of ϕ to maximize $\|\mathbf{X}\phi\|_2^2$, *s.t.* $\|\phi\|_2 = 1$ is exactly the first column of \mathbf{U} , where $[\mathbf{U}, \mathbf{S}] = \text{svd}(\mathbf{X}^T \mathbf{X})$. (Note: you need prove why it is optimal than any other reasonable combinations of \mathbf{U}_i , say $\hat{\phi} = 0.8 * \mathbf{U}(:, 1) + 0.6 * \mathbf{U}(:, 2)$ which also satisfies $\|\hat{\phi}\|_2 = 1$.)

Problem 2

Given matrix $\mathbf{X} \in \mathbb{R}^{n \times p}$ (assume each column is centered already), where n denotes sample size while p feature size. To conduct PCA, we need find eigenvectors to the largest eigenvalues of $\mathbf{X}^T \mathbf{X}$, where usually the complexity is $\mathcal{O}(p^3)$. Apparently when $n \ll p$, this is not economic when p is large. Please consider conducting PCA based on $\mathbf{X}\mathbf{X}^T$ and obtain the eigenvectors for $\mathbf{X}^T \mathbf{X}$ accordingly and use experiment to demonstrate the acceleration.

Problem 3

Let's revisit Least Squares Problem: minimize $\frac{1}{2}\|\mathbf{y} - \mathbf{A}\boldsymbol{\beta}\|_2^2$, where $\mathbf{A} \in \mathbb{R}^{n \times p}$.

1. Please show that if $p > n$, then vanilla solution $(\mathbf{A}^T \mathbf{A})^{-1} \mathbf{A}^T \mathbf{y}$ is not applicable any more.
2. Let's assume $\mathbf{A} = [1, 2, 4; 1, 3, 5; 1, 7, 7; 1, 8, 9]$, $\mathbf{y} = [1; 2; 3; 4]$. Please show via experiment results that Gradient Descent method will obtain the optimal solution with Linear Convergence rate if the learning rate is fixed to be $\frac{1}{\sigma_{\max}(\mathbf{A}^T \mathbf{A})}$, and $\boldsymbol{\beta}_0 = [0; 0; 0]$.
3. Now let's consider ridge regression: minimize $\frac{1}{2}\|\mathbf{y} - \mathbf{A}\boldsymbol{\beta}\|_2^2 + \frac{\lambda}{2}\|\boldsymbol{\beta}\|_2^2$, where $\mathbf{A}, \mathbf{y}, \boldsymbol{\beta}_0$ remains the same as above while learning rate is fixed to be $\frac{1}{\lambda + \sigma_{\max}(\mathbf{A}^T \mathbf{A})}$ where λ varies from 0.1, 1, 10, 100, 200, please show that Gradient Descent method with larger λ converges faster.

Problem 4

We consider matrix completion problem. As we discussed in class, the main issue of *softImpute* (*Matrix Completion via Iterative Soft-Thresholded SVD*) is when the matrix size is large, conducting SVD is computational demanding. Let's recall the original problem where $\mathbf{X}, \mathbf{Z} \in \mathbb{R}^{n \times d}$:

$$\min_{\mathbf{Z}} \frac{1}{2} \|P_{\Omega}(\mathbf{X}) - P_{\Omega}(\mathbf{Z})\|_F^2 + \lambda \|\mathbf{Z}\|_* \quad (1)$$

People have found that instead of finding optimal \mathbf{Z} , it might be better to make use of *Burer-Monteiro* method to optimize two matrices $\mathbf{A} \in \mathbb{R}^{n \times r}$, $\mathbf{B} \in \mathbb{R}^{d \times r}$ ($r \geq \text{rank}(\mathbf{Z}^*)$) such that $\mathbf{A}\mathbf{B}^T = \mathbf{Z}$. The new objective is:

$$\min_{\mathbf{A}, \mathbf{B}} \frac{1}{2} \|P_{\Omega}(\mathbf{X} - \mathbf{A}\mathbf{B}^T)\|_F^2 + \frac{\lambda}{2} (\|\mathbf{A}\|_F^2 + \|\mathbf{B}\|_F^2). \quad (2)$$

- Assume $[\mathbf{U}, \boldsymbol{\Sigma}, \mathbf{V}] = \text{svd}(\mathbf{Z})$, show that if $\mathbf{A} = \mathbf{U}\boldsymbol{\Sigma}^{\frac{1}{2}}$, $\mathbf{B} = \mathbf{V}\boldsymbol{\Sigma}^{\frac{1}{2}}$, then Eq. (2) is equivalent to Eq. (1).
- The *Burer-Monteiro* method suggests if we can find $\mathbf{A}^*, \mathbf{B}^*$, then the optimal \mathbf{Z} to Eq. (1) can be recovered by $\mathbf{A}^* \mathbf{B}^{*T}$. It boils down to solve Eq. (2). Show that we can make use of least squares with ridge regression to update \mathbf{A}, \mathbf{B} row by row in an alternating minimization manner as below. Assume $n = d = 2000$, $r = 200$, please write program to find \mathbf{Z}^* .

$T \leftarrow 100, i \leftarrow 1$ % you can also set T to be other number instead of 100

if $i \leq T$ **then**

 update \mathbf{A} row by row while fixing \mathbf{B}

 update \mathbf{B} row by row while fixing \mathbf{A}

$i \leftarrow i + 1$

end if

Ans 1.

For matrix with self centered data,
max. variance is given by

$$\max \|\mathbf{X}\phi\|_2^2 \quad \text{---} \quad (1)$$

such that $\|\phi\|_2 = 1$ and ϕ is a vector

To prove,

Simplest closed solution of (1) is exactly
equal to the first column of \mathbf{U}

where $(\mathbf{U}, \mathbf{\Sigma}) = \text{svd}(\mathbf{X}^T \mathbf{X})$

modify $\|\mathbf{X}\phi\|_2^2$ as

$$[\mathbf{X}\phi]^T [\mathbf{X}\phi] = \phi^T \mathbf{X}^T \mathbf{X} \phi$$

As $\mathbf{X}^T \mathbf{X}$ is a spd matrix,
for SVD of $\mathbf{X}^T \mathbf{X}$, we have

$$\mathbf{U} = \mathbf{U} \quad \text{in } \mathbf{U} \mathbf{\Sigma} \mathbf{V}^T$$

Also, singular values are absolute values of
eigen values

Also, columns of $U = V$ are eigenvectors of A

Consider:

Case I - when $\phi = aV_1 + bV_2$

where $a, b \in \mathbb{R}$

s.t. $a^2 + b^2 = 1$

So,

$$\begin{aligned} \|\phi\|^2 &= \phi^T \cdot \phi \\ &= (aV_1 + bV_2)^T (aV_1 + bV_2) \\ &= a^2 V_1^T V_1 + b^2 V_2^T V_2 + 2ab V_1^T V_2 \end{aligned}$$

$$\begin{aligned} (\because V_1 \perp V_2) \\ V_1^T V_1 = 1 \\ V_2^T V_2 = 1 \end{aligned}$$

So,

$$\|\phi\|^2 = a^2 + b^2 = 1$$

s.t. $a^2 + b^2 = 1$

satisfy the given constraint

If we substitute the value of ϕ in $\phi^T X^T X \phi$

$$\begin{aligned} &= \phi^T [U \Sigma U^T] \phi \\ &= \phi^T [\sigma_1 V_1 V_1^T + \sigma_2 V_2 V_2^T + \sigma_3 V_3 V_3^T + \dots] \phi \end{aligned}$$

$$= (aV_1 + bV_2)^T [\sigma_1 V_1 V_1^T + \sigma_2 V_2 V_2^T + \sigma_3 V_3 V_3^T + \dots] (aV_1 + bV_2)$$

$$= (aV_1^T \sigma_1 V_1 V_1^T + bV_2^T \sigma_2 V_2 V_2^T + \dots) (aV_1 + bV_2)$$

$$= (a^2 \sigma_1 V_1^T V_1 V_1^T V_1 + b^2 \sigma_2 V_2^T V_2 V_2^T V_2 + \dots)$$

As $V_1 \perp V_2$

$$V_1^T V_2 = V_2^T V_1 = 0$$

So, we are left with

$$a^2 \sigma_1 + b^2 \sigma_2$$

$$\sigma_1 = \left(a^2 + b^2 \frac{\sigma_2^2}{\sigma_1^2} \right) \quad \text{--- (1)}$$

As we know that the largest singular value is σ_1 & it is always greater than σ_2 and,

$$a^2 + b^2 = 1$$

Therefore, we have

$$a^2 + b^2 \left(\frac{\sigma_2}{\sigma_1} \right) < 1$$

$$\sigma_1 \left(a^2 + b^2 \left(\frac{\sigma_2}{\sigma_1} \right) \right) < \sigma_1 \quad - (2)$$

Case II consider if $\phi = v_1$

$$\|\phi\|_2^2 = 1 \quad (\because \text{orthogonal vector})$$

Put $\phi = v_1$ in

$$\phi^T [U \Sigma U^T] \phi$$

$$\begin{aligned} & \phi^T [\sigma_1 v_1 v_1^T + \sigma_2 v_2 v_2^T + \sigma_3 v_3 v_3^T + \dots] \phi \\ &= \sigma_1 v_1^T v_1 v_1^T v_1 + \sigma_2 v_2^T v_1 v_2^T v_1 + \dots \end{aligned}$$

As all columns of U are orthogonal,

we have,

$$\sigma_1 v_1^T v_1 v_1^T v_1$$

$$\text{Also } v_1^T v_1 = I$$

$$\max \|X\phi\|_2^2 = \sigma_1$$

Comparing $w = v$ as an argument to a general combination of vectors from V ,

eq. 3 gives:

$$\text{arg max}_{\|\phi\|_2^2=1} \|X\phi\|_2^2 \Rightarrow w = aV_1 + bV_2$$

(s.t. $a, b \in \mathbb{R}$ & $a^2 + b^2 = 1$)

$$\text{gives } \sigma_1 (a^2 + b^2 \frac{\sigma_2}{\sigma_1})$$

which is less than σ_1 when $w \neq v_1$

So, max variance is achieved when $\phi = \text{first column of } V$, where

$$V = [V, S] = \text{SVD}(X^T X)$$

Ans 2

$$\text{let } A = X^T X \\ B = X X^T$$

let U be eigen vectors of B

$$X X^T U = \lambda U$$

multiply both sides by X^T

$$X^T X X^T U = \lambda X^T U$$

from ①, $X^T X = A$

$$A(X^T U) = \lambda(X^T U)$$

So, we can conclude that $X^T U$ are eigen vectors of A .

So, we can find eigen vectors of B , which have the computational complexity of $O(N^3)$ instead of $O(p^3)$, when $X \in \mathbb{R}^{n \times p}$

These eigen vectors can be used to find eigen vectors of $A(X^T X)$

To convert eigen vectors of xx^T to those of $x^T x$

→ Multiply eigen vectors of $B(xx^T)$ by x^T and then normalize

i.e. eigen vectors of $A: x^T U$

But $x^T U$ is not normalized. So, the matrix which gives us eigen vectors of A is $x^T U$ with all of its column vectors normalized.

Problem 2

```
% Experimental Setup
samples = 100;
features = 10000;

% Generate random data for matrix X (centered)
MatrixX = randn(samples, features);

% PCA using  $X^T X$ 
start_XTX = tic;
CovMatrix1 = MatrixX' * MatrixX;
[EigVecs1, EigVals1] = eig(CovMatrix1);
[~, order1] = sort(diag(EigVals1), 'descend');
EigVecs1 = EigVecs1(:, order1);
elapsed_XTX = toc(start_XTX);

% PCA using  $XX^T$ 
start_XXT = tic;
CovMatrix2 = MatrixX * MatrixX';
[EigVecs2, EigVals2] = eig(CovMatrix2);
EigVecs2 = MatrixX' * EigVecs2;
for k = 1:samples
    EigVecs2(:,k) = EigVecs2(:,k) / norm(EigVecs2(:,k));
end
[~, order2] = sort(diag(EigVals2), 'descend');
EigVecs2 = EigVecs2(:, order2);
elapsed_XXT = toc(start_XXT);

% Display the results
fprintf('Time for PCA using  $X^T X$ : %.4f seconds\n', elapsed_XTX);
```

Time for PCA using $X^T X$: 63.1824 seconds

```
fprintf('Time for PCA using  $XX^T$ : %.4f seconds\n', elapsed_XXT);
```

Time for PCA using XX^T : 0.0580 seconds

Problem 3.1

Case: $n < p$

```
n = 10; % number of samples
p = 10000; % number of features

% Generate random data for matrix X (centered)
X = randn(n, p);

Y = X'* X;
disp("If n < p:")
```

If $n < p$:

```
disp("Size of the square matrix:");
```

Size of the square matrix:

```
disp(size(Y));
```

10000 10000

```
disp("Rank of the square matrix");
```

Rank of the square matrix

```
disp(rank(Y));
```

10

Case $n > p$:

```
disp("If n > p:");
```

If $n > p$:

```
n = 1000; % number of samples
p = 100; % number of features

% Generate random data for matrix X (centered)
X = randn(n, p);

Y = X'*X;
disp("Size of the square matrix:");
```

Size of the square matrix:

```
disp(size(Y));
```

100 100

```
disp("Rank of the square matrix");
```

Rank of the square matrix

```
disp(rank(Y));
```

100

Given that the rank of the matrix is less than its size, and the number of features p exceeds the number of samples n , the resulting matrix is singular and therefore non-invertible. Consequently, this scenario precludes the viability of a closed-form solution.

PROBLEM 3.2

```
% Given matrix A, target vector y, learning rate, and initial weights
```

```
A = [1, 2, 4; 1, 3, 5; 1, 7, 7; 1, 8, 9];
```

```
y = [1; 2; 3; 4];
```

```
% Compute the maximum singular value of A
```

```
max_singular_value = max(svd(A' * A));
```

```
% Set learning rate as the inverse of the max singular value
```

```
learningRate = 1 / max_singular_value;
```

```
weights = zeros(size(A, 2), 1);
```

```
% Initialize a vector to track loss values
```

```
lossHistory = zeros(1, 32000);
```

```
% Vanilla Gradient Descent (32k iterations)
```

```
for epoch = 1:32000
```

```
    gradient = A' * (A * weights - y);
```

```
    weights = weights - learningRate * gradient;
```

```
    lossHistory(epoch) = norm(A * weights - y)^2;
```

```
end
```

```
% Calculate the convergence factor
```

```
convergenceFactor = lossHistory(2:end) ./ lossHistory(1:end-1);
```


```
% Plot Loss vs Epochs
```

```
plot(lossHistory);
```

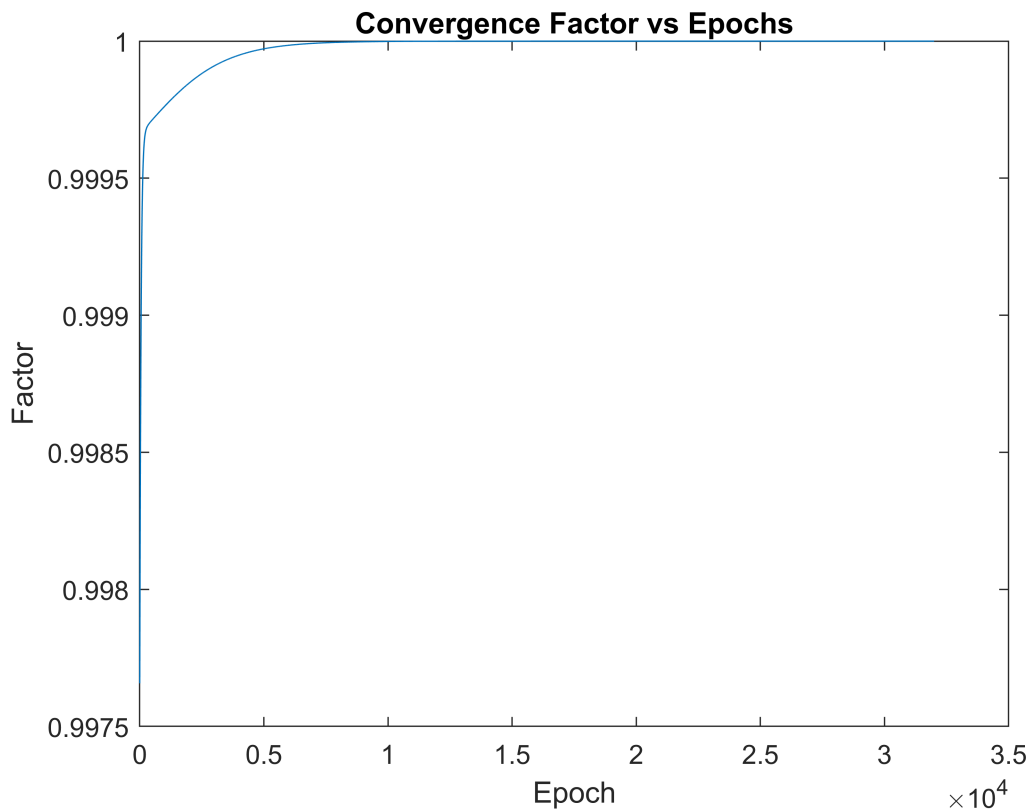
```
title('Loss vs Epochs');
```

```
xlabel('Epochs');
```

```
ylabel('Loss');
```

```
% Plot Convergence Factor vs Epochs
figure;
plot(convergenceFactor);
title('Convergence Factor vs Epochs');
xlabel('Epoch');
ylabel('Factor');
```



```
% Closed form solution for verification
weightsClosedForm = pinv(A' * A) * A' * y;
display(weightsClosedForm);
```

```
weightsClosedForm = 3x1
-1.0000
 0.0333
 0.5333
```

```
display(weights);
```

```
weights = 3x1
-1.0000
 0.0333
 0.5333
```

From analyzing the weight matrices, it's evident that the Gradient Descent approach achieves the optimal solution, exhibiting nearly linear convergence with the specified learning rate.

PROBLEM 3.3

```
% Define the matrix A and vector y
A = [1, 2, 4; 1, 3, 5; 1, 7, 7; 1, 8, 9];
y = [1; 2; 3; 4];
```

```
% Define a range of  $\lambda$  values
lambda_values = [0.1, 1, 10, 100, 200];
```



```

% Define the number of iterations
num_iterations = 20;

% Initialize variables to store losses
losses = zeros(length(lambda_values), num_iterations);

% Perform gradient descent for each  $\lambda$ 
for k = 1:length(lambda_values)
    lambda = lambda_values(k);
    alpha = 1 / (max(max(A' * A)) + lambda_values(k));

    % Initialize  $\theta$  to  $[0; 0; 0]$ 
    theta = [0; 0; 0];
    for iteration = 1:num_iterations
        % Compute the gradient for ridge regression
        gradient = -A' * (y - A * theta) + lambda * theta;


        % Update  $\theta$  using the gradient and learning rate
        theta = theta - alpha * gradient;

        % Calculate the cost (ridge regression loss)
        loss = (1 / (2)) * norm(A * theta - y)^2 + (theta' ...
            * theta) * (lambda / 2);
        % Store the loss
        losses(k, iteration) = loss;
    end
end

% Plot the loss curves for different  $\lambda$  values
figure;
hold on;
for k = 1:length(lambda_values)
    semilogy(1:num_iterations, losses(k, :), 'DisplayName', ...
        ['\lambda = ', num2str(lambda_values(k))]);
end

hold off;
xlabel('Iteration');
ylabel('Loss');
title("Loss vs Epochs");
legend;

```



%Verifying the results obtained from closed form solution

```
weights2 = pinv(A'*A + lambda_values(5)*eye(size(A, 2)))*A'*y;
```

To verify, let us compare the closed form weights and weights

```
display(weights2);
```

```
weights2 = 3x1
    0.0195
    0.1231
    0.1423
```

```
display(weights);
```

```
weights = 3x1
   -1.0000
    0.0333
    0.5333
```


Ans 4.

$$\min_{A, B} \frac{1}{2} \|P_{\Omega}(X - AB^T)\|_F^2 + \frac{\lambda}{2} (\|A\|_F^2 + \|B\|_F^2) \quad \text{--- (1)}$$

Given : $A = U \Sigma^{1/2} \quad B = V \Sigma^{1/2}$

$$[U, \Sigma, V] = \text{SVD}(Z)$$

$$\begin{aligned} P_{\Omega}(X - AB^T) &= P_{\Omega}(X) - P_{\Omega}(AB^T) \\ &= P_{\Omega}(X) - P_{\Omega}(Z) \quad \text{--- (2)} \end{aligned}$$

$$\begin{aligned} \|A\|_F^2 &= \text{tr}(A^T A) = \text{tr}[(U \Sigma^{1/2})^T (U \Sigma^{1/2})] \\ &= \text{tr}[(\Sigma^{1/2})^T U^T \cdot U \Sigma^{1/2}] \end{aligned}$$

$$A \quad U^T \cdot U = I$$

$$\|A\|_F^2 = \text{tr}[(\Sigma^{1/2})^T (\Sigma^{1/2})] = \text{tr}(\Sigma), \quad \text{--- (3)}$$

Similarly, $B^T B = (V \Sigma^{1/2})^T (V \Sigma^{1/2})$

$$\|B\|_F^2 = \text{tr}(B^T B) = \text{tr}(\Sigma), \quad \text{--- (4)}$$

Putting (2), (3), (4) in (1)

$$\min_{A, B} \frac{1}{2} \|P_R(X) - P_R(Z)\|_F^2 + \frac{\lambda}{2} (\text{tr}(\varepsilon) + \text{tr}(\varepsilon))$$

$$= \min_{A, B} \frac{1}{2} \|P_R(X) - P_R(Z)\|_F^2 + \lambda \text{tr}(\varepsilon)$$

$\text{tr}(\varepsilon)$ is sum of ~~the~~ the singular values

\therefore it is equal to nuclear norm of Z

Therefore, we have

$$\min_{A, B} \frac{1}{2} \|P_R(X) - P_R(Z)\|_F^2 + \lambda \|Z\|_*$$

PROBLEM 4.2

```
% Parameters
n = 2000; % Number of samples
d = 2000; % Number of features
r = 200;  % Rank for A and B

% Generate a random matrix X for testing
X = randn(n, d);

% Regularization parameter
lambda = 0.1;

% Initialize matrices A and B
A = randn(n, r);
B = randn(d, r);

% Maximum number of iterations
num_iterations = 100;

% Learning rate
learning_rate = 1 / (max(svd(A))^2 + lambda);

% Iterative updates for A and B
for iter = 1:num_iterations
    AB_diff = X - A * B';
    A = A + learning_rate * (AB_diff * B - lambda * A);
    B = B + learning_rate * (AB_diff' * A - lambda * B);
end
```